Assignment: IDfy

• @Name : Subham Sarkar

• @Github: https://github.com/SubhamIO

• @LinkedIn: https://www.linkedin.com/in/subham-sarkar-4224aa147/ • @Portfolio: https://subhamio.github.io/SubhamSarkar-PortfolioWebsite/

Title: Text Recognition from low quality License Plates using **Deep Learning**

Objective:

• The objective of this assignment is to build an OCR solution for the provided dataset. This specific dataset is normal and HDR readings of license plates.

Requirements:

- 1. Use an 80:20 train:test split on the provided dataset
- 2. Create a model for reading the text using any approaches or tools that you are familiar with or can learn
- 3. Use the available test set to check the accuracy of your model

Dataset:

- 1. The image dataset here contains 652 images of cropped license plates with a csv containing annotation as well.
- Link: https://medusa.fit.vutbr.cz/traffic/download/513/

Load Libraries

```
In [1]:
```

import cv2

import numpy as np

from matplotlib import pyplot as plt

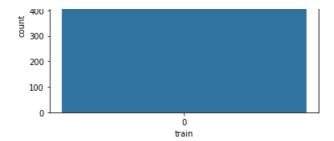
```
from google.colab import drive
drive.mount('/content/drive')
Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client id=947318989803-6bn6
qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect uri=urn%3aietf%3awg%3aoauth%3a2.0%
b&response type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2
www.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly
ttps%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly
Enter your authorization code:
Mounted at /content/drive
4
In [2]:
cd drive/My\ Drive/
/content/drive/My Drive
In [3]:
import re
```

```
import pandas as pd
from keras.preprocessing.image import *
from keras.layers.core import *
import tensorflow as tf
from keras.layers import *
from keras.models import *
import keras
from keras import backend as K
from keras.callbacks import *
np.random.seed(0)
from keras.utils.np_utils import to_categorical
from keras.regularizers import 12
import seaborn as sns
Using TensorFlow backend.
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/ testing.py:19: FutureWarning:
pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.
  import pandas.util.testing as tm
Data Acquisition
In [0]:
# base dir = '/Users/subham/Desktop/2017-IWT4S-HDR LP-dataset/'
base dir = "./2017-IWT4S-HDR LP-dataset/"
In [0]:
data = pd.read_csv("./2017-IWT4S-HDR_LP-dataset/" + "/trainVal.csv")
In [6]:
data.tail()
Out[6]:
              track id
                            image_path
                                           lp train
647 ./crop_m4/I00084.png
                      ./crop_h4/I00084.png 2B90178
648 ./crop_m4/I00085.png
                     ./crop_m4/I00085.png 7B59839
649 ./crop_m4/I00085.png
                      ./crop_h4/I00085.png 7B59839
                                                0
650 ./crop_m4/I00086.png
                     ./crop_m4/I00086.png 7B11123
651 ./crop_m4/I00086.png
                      ./crop_h4/I00086.png 7B11123
In [17]:
n = data.shape[0]
print(n)
652
ax = sns.countplot(x=data['train'], data=data)
print(data['train'].value counts())
```

600 -500 -

652

Name: train, dtype: int64



Observation:

• It has only test data as per README.txt provided

```
In [19]:
len(data['track_id'].unique())
Out[19]:
326

In [20]:
len(data['image_path'].unique())
Out[20]:
652

In [21]:
len(data['lp'].unique())
Out[21]:
302
```

Observation:

- Unique track id = 326
- Unique image_path = 652
- Many : 1 relationship [track_id : image_path]
- Unique lp = 302, this means same images are present in different folders. So, not much variation in image types also the dataset is small. Accuracy will be impacted.

Data Preprocessing

Create a dictionary of length 37 [26 alphabets and 10 numbers and space] for target map

```
In [0]:

letters = " ABCDEFGHIJKLMNPQRSTUVWXYZ0123456789"
dic = {}
for i in range(len(letters)):
    dic[i] = letters[i]
invert_dic = {}
for i in range(len(letters)):
    invert_dic[letters[i]] = i
```

```
print(dic)
{0: ' ', 1: 'A', 2: 'B', 3: 'C', 4: 'D', 5: 'E', 6: 'F', 7: 'G', 8: 'H', 9: 'I', 10: 'J', 11: 'K',
12: 'L', 13: 'M', 14: 'N', 15: 'P', 16: 'Q', 17: 'R', 18: 'S', 19: 'T', 20: 'U', 21: 'V', 22: 'W',
23: 'X', 24: 'Y', 25: 'Z', 26: '0', 27: '1', 28: '2', 29: '3', 30: '4', 31: '5', 32: '6', 33: '7',
34: '8', 35: '9'}
In [24]:
print(invert dic)
{' ': 0, 'A': 1, 'B': 2, 'C': 3, 'D': 4, 'E': 5, 'F': 6, 'G': 7, 'H': 8, 'I': 9, 'J': 10, 'K': 11,
'L': 12, 'M': 13, 'N': 14, 'P': 15, 'Q': 16, 'R': 17, 'S': 18, 'T': 19, 'U': 20, 'V': 21, 'W': 22, 'X': 23, 'Y': 24, 'Z': 25, '0': 26, '1': 27, '2': 28, '3': 29, '4': 30, '5': 31, '6': 32, '7': 33,
'8': 34, '9': 35}
In [0]:
X train = []
y_train = []
X \text{ test} = []
y test = []
Creating X(feature), Y(target)
In [0]:
X = []
Y = []
In [0]:
for i in range(n):
    temp_y= np.zeros((8)) # Target variable size = 8 dimensional
    path = base dir + data["image path"][i]
     # Read the images in gray scale
    temp x = cv2.imread(base dir + data["image path"][i], cv2.IMREAD GRAYSCALE)
    # Resizing as per our need to process in our CNN architecture
    temp_x = cv2.resize(temp_x, (256, 64))
    X.append(temp_x)
    # Let's loop over each ground truth and assign each character with index
    for j,k in enumerate(data["lp"][i]):
        temp y[j] = invert dic[k]
    Y.append(temp_y)
In [28]:
len(X), len(Y)
Out[28]:
(652, 652)
In [0]:
# import pickle
# with open('X_train', 'wb') as fp:
    pickle.dump(X train, fp)
In [0]:
# with open('X test', 'wb') as fp:
     pickle.dump(X_test, fp)
```

```
In [0]:
# with open('y train', 'wb') as fp:
     pickle.dump(y train, fp)
In [0]:
# with open('y test', 'wb') as fp:
     pickle.dump(y_test, fp)
In [0]:
# with open ('y test', 'rb') as fp:
     y_test1 = pickle.load(fp)
Let's look at the mapping created
In [30]:
data['lp'][0]
Out[30]:
'9B52145'
{' ': 0, 'A': 1, 'B': 2, 'C': 3, 'D': 4, 'E': 5, 'F': 6, 'G': 7, 'H': 8, 'I': 9, 'J': 10, 'K': 11, 'L': 12, 'M': 13, 'N': 14, 'P': 15, 'Q': 16, 'R': 17, 'S':
18, 'T': 19, 'U': 20, 'V': 21, 'W': 22, 'X': 23, 'Y': 24, 'Z': 25, '0': 26, '1': 27, '2': 28, '3': 29, '4': 30, '5': 31, '6': 32, '7': 33, '8': 34, '9': 35}
In [31]:
Y[0]
Out[31]:
array([35., 2., 31., 28., 27., 30., 31., 0.])
Train-Test Split (80:20)
In [32]:
11 = 0.8*len(X)
ll = int(ll)
X train = X[:11]
X test = X[ll:]
y_train = Y[:11]
y test = Y[ll:]
print(len(X train), len(y train), len(X test), len(y test))
521 521 131 131
In [0]:
# reshaping the array [-1 automatically adjust the number of data points here]
X \text{ train} = \text{np.array}(X \text{ train}).\text{reshape}(-1,64,256,1)
y train = np.array(y train)
X_{\text{test}} = \text{np.array}(X_{\text{test}}).\text{reshape}(-1,64,256,1)
y_test = np.array(y_test)
# Normalise the data
X train = X train /255
X_{test} = X_{test/255}
# converting target to categorical (encoding the target characters to 36 dimensions)
y_test = to_categorical(y_test,36)
v train = to categorical(v train.36)
```

```
In [34]:
print(X train.shape)
print(y train.shape)
print(X test.shape)
print(y_test.shape)
(521, 64, 256, 1)
(521, 8, 36)
(131, 64, 256, 1)
(131, 8, 36)
In [0]:
print(y train[0])
0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0.]
0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]
0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]
0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]]
```

Data Modelling

Let's define VGG-16 architecture

- Could have used transfer learning but VGG-16 takes (224,224,3) as input dimension
- So, better define the VGG-16 architecture from scratch, so that we can specify input dimensions manually.

```
def VGG(shape=(64, 256, 1),n_channels=64,weight_decay=0,batch_momentum=0.99):
                bn axis = 3
                 input_ = Input (shape=shape)
                 x = Conv2D(128, (3, 3), padding='same', name='block1_conv1', kernel_initializer='he_normal', ke
rnel_regularizer=12(weight_decay))(input_)
                 x = BatchNormalization(axis=bn axis, name='bn00 x1', momentum=batch momentum)(x)
                 x = Activation('relu')(x)
                 x = Conv2D(128, (3, 3), padding='same', name='block1_conv2', kernel_initializer='he_normal', kernel_initializer='he_normal',
rnel regularizer=12(weight decay))(x)
                 x = BatchNormalization(axis=bn axis, name='bn01 x2', momentum=batch momentum)(x)
                 x = Activation('relu')(x)
                 x = MaxPooling2D((2, 2), strides=(2, 2), name='block1 pool')(x)
                 # Block 2
                 x = Conv2D(128, (3, 3), padding='same', name='block2 conv1', kernel initializer='he normal', ke
rnel regularizer=12(weight decay))(x)
                x = BatchNormalization(axis=bn axis, name='bn11 x1', momentum=batch momentum)(x)
                 x = Activation('relu')(x)
                 x = Conv2D(128, (3, 3), padding='same', name='block2\_conv2', kernel\_initializer='he_normal', kernel\_initializer='he_normal',
rnel regularizer=12(weight decay))(x)
                 x = BatchNormalization(axis=bn axis, name='bn12 x2', momentum=batch momentum)(x)
                 x = Activation('relu')(x)
                 x = MaxPooling2D((2, 2), strides=(2, 2), name='block2 pool')(x)
                 # Block 3
```

```
x = Conv2D(256, (3, 3), padding='same', name='block3 conv1', kernel initializer='he normal', ke
rnel regularizer=12(weight decay))(x)
   x = BatchNormalization(axis=bn_axis, name='bn21_x1', momentum=batch momentum)(x)
   x = Activation('relu')(x)
   x = Conv2D(256, (3, 3), padding='same', name='block3_conv2', kernel_initializer='he_normal', ke
rnel regularizer=12(weight decay))(x)
   x = BatchNormalization(axis=bn axis, name='bn22 x2', momentum=batch momentum)(x)
   x = Activation('relu')(x)
   x = Conv2D(256, (3, 3), padding='same', name='block3 conv3',
kernel_initializer='glorot_uniform', kernel_regularizer=12(weight_decay))(x)
   x = BatchNormalization(axis=bn_axis, name='bn23_x3', momentum=batch_momentum)(x)
    x = Activation('relu')(x)
   x = MaxPooling2D((2, 2), strides=(2, 2), name='block3 pool')(x)
   x = Conv2D(512, (3, 3), padding='same', name='block4_conv1', kernel_initializer='he_normal', ke
rnel regularizer=12(weight_decay))(x)
   x = BatchNormalization(axis=bn axis, name='bn31 x2', momentum=batch momentum)(x)
   x = Activation('relu')(x)
   x = Conv2D(512, (3, 3), padding='same', name='block4 conv2', kernel regularizer=12(weight decay
))(x)
   x = BatchNormalization(axis=bn axis, name='bn32 x2', momentum=batch momentum)(x)
   x = Activation('relu')(x)
    x = Conv2D(512, (3, 3), padding='same', name='block4\_conv3',
kernel initializer='glorot uniform', kernel regularizer=12(weight decay))(x)
   x = BatchNormalization(axis=bn_axis, name='bn33_x2', momentum=batch_momentum)(x)
    x = Activation('relu')(x)
   x = MaxPooling2D((2, 2), strides=(2, 2), name='block4 pool')(x)
   # Block 5
   x = Conv2D(512, (3, 3), padding='same', name='block5 conv1', kernel initializer='he normal', ke
rnel_regularizer=12(weight_decay))(x)
   x = BatchNormalization(axis=bn axis, name='bn41 x2', momentum=batch momentum)(x)
   x = Activation('relu')(x)
    x = Conv2D(512, (3, 3), padding='same', name='block5 conv2',
kernel initializer='qlorot uniform', kernel regularizer=12(weight decay))(x)
   x = BatchNormalization(axis=bn axis, name='bn42 x2', momentum=batch momentum)(x)
    x = Activation('relu')(x)
   x = Conv2D(1024, (3, 3), padding='same', name='block5_conv3', kernel_initializer='he_normal', k
ernel regularizer=12(weight decay))(x)
   x = BatchNormalization(axis=bn_axis, name='bn43_x2', momentum=batch_momentum)(x)
   x = Activation('relu')(x)
   x = MaxPooling2D((2, 2), strides=(2, 2), name='block5 pool')(x)
   x = Conv2D(1024, (3, 3), padding='same', name='block6 conv1', kernel initializer='he normal', k
ernel regularizer=12(weight_decay))(x)
   x = BatchNormalization(axis=bn axis, name='bn51 x2', momentum=batch momentum)(x)
   x = Activation('relu')(x)
   x = Conv2D(1024*2, (3, 3), padding='same', name='block6 conv12', kernel initializer='he normal'
, kernel regularizer=12(weight decay))(x)
   x = BatchNormalization(axis=bn axis, name='bn51 x22', momentum=batch momentum)(x)
   x = Activation('relu')(x)
   x= Dropout(0.3, noise shape=None, seed=None)(x)
   #block5
   X = AveragePooling2D((2, 2), strides = (2, 1), name='avg pool1', padding ='same')(x)
   X = Reshape((8,1024*2))(X)
   X = Conv1D(512, 3, strides=1, padding='same', name = 'conv1y' ,activation=None, dilation_rate=
1, use bias=True, kernel initializer="he normal", kernel regularizer=regularizers.12(0.000))(X)
   X = BatchNormalization(axis = 2, name = 'bn01y')(X)
   X = Activation('relu')(X)
   X= Dropout(0.3, noise shape=None, seed=None)(X)
   X = Conv1D(36, 1 , strides=1, padding='same', name = 'conv1x' ,activation=None, dilation_rate=
1, use bias=True, kernel initializer="he normal")(X)
   X = BatchNormalization(axis = 2, name = 'bnhe')(X)
   X = Activation('softmax')(X)
   model = Model(inputs = [input], outputs = [X])
   return model
```

In [0]:

```
model = VGG(shape=(64, 256, 1))
```

In [0]:

model.summary()

Model: "model_1"

Layer (type)	Output Sl	-	Param #
input_1 (InputLayer)	(None, 6	4, 256, 1)	0
block1_conv1 (Conv2D)	(None, 6	4, 256, 128)	1280
bn00_x1 (BatchNormalization)	(None, 6	4, 256, 128)	512
activation_1 (Activation)	(None, 6	4, 256, 128)	0
block1_conv2 (Conv2D)	(None, 6	4, 256, 128)	147584
bn01_x2 (BatchNormalization)	(None, 6	4, 256, 128)	512
activation_2 (Activation)	(None, 6	4, 256, 128)	0
block1_pool (MaxPooling2D)	(None, 32	2, 128, 128)	0
block2_conv1 (Conv2D)	(None, 32	2, 128, 128)	147584
bn11_x1 (BatchNormalization)	(None, 32	2, 128, 128)	512
activation_3 (Activation)	(None, 32	2, 128, 128)	0
block2_conv2 (Conv2D)	(None, 32	2, 128, 128)	147584
bn12_x2 (BatchNormalization)	(None, 32	2, 128, 128)	512
activation_4 (Activation)	(None, 32	2, 128, 128)	0
block2_pool (MaxPooling2D)	(None, 1	6, 64, 128)	0
block3_conv1 (Conv2D)	(None, 1	6, 64, 256)	295168
bn21_x1 (BatchNormalization)	(None, 1	6, 64, 256)	1024
activation_5 (Activation)	(None, 1	6, 64, 256)	0
block3_conv2 (Conv2D)	(None, 1	6, 64, 256)	590080
bn22_x2 (BatchNormalization)	(None, 1	6, 64, 256)	1024
activation_6 (Activation)	(None, 1	6, 64, 256)	0
block3_pool (MaxPooling2D)	(None, 8	, 32, 256)	0
block4_conv1 (Conv2D)	(None, 8	, 32, 512)	1180160
bn31_x2 (BatchNormalization)	(None, 8	, 32, 512)	2048
activation_7 (Activation)	(None, 8	, 32, 512)	0
block4_conv2 (Conv2D)	(None, 8	, 32, 512)	2359808
bn32_x2 (BatchNormalization)	(None, 8	, 32, 512)	2048
activation_8 (Activation)	(None, 8	, 32, 512)	0
block4_pool (MaxPooling2D)	(None, 4	, 16, 512)	0
block5_conv1 (Conv2D)	(None, 4	, 16, 512)	2359808
bn41_x2 (BatchNormalization)	(None, 4	, 16, 512)	2048
			-

activation_9 (Activation)	(None,	4,	16, 512)	0
block5_conv3 (Conv2D)	(None,	4,	16, 1024)	4719616
bn43_x2 (BatchNormalization)	(None,	4,	16, 1024)	4096
activation_10 (Activation)	(None,	4,	16, 1024)	0
block5_pool (MaxPooling2D)	(None,	2,	8, 1024)	0
block6_conv1 (Conv2D)	(None,	2,	8, 1024)	9438208
bn51_x2 (BatchNormalization)	(None,	2,	8, 1024)	4096
activation_11 (Activation)	(None,	2,	8, 1024)	0
block6_conv12 (Conv2D)	(None,	2,	8, 2048)	18876416
bn51_x22 (BatchNormalization	(None,	2,	8, 2048)	8192
activation_12 (Activation)	(None,	2,	8, 2048)	0
dropout_1 (Dropout)	(None,	2,	8, 2048)	0
<pre>avg_pool1 (AveragePooling2D)</pre>	(None,	1,	8, 2048)	0
reshape_1 (Reshape)	(None,	8,	2048)	0
convly (ConvlD)	(None,	8,	512)	3146240
bn01y (BatchNormalization)	(None,	8,	512)	2048
activation_13 (Activation)	(None,	8,	512)	0
dropout_2 (Dropout)	(None,	8,	512)	0
convlx (Conv1D)	(None,	8,	36)	18468
bnhe (BatchNormalization)	(None,	8,	36)	144
activation_14 (Activation)	(None,	8,	36)	0
Total params: 43,456,820 Trainable params: 43,442,412 Non-trainable params: 14,408				

Defining custom metrics

```
In [0]:
```

```
def custom_loss(y_true, y_pred):
    s = K.shape(y_pred)
    y_true = K.reshape(y_true, (-1,s[-1]))
    y_pred = K.reshape(y_pred, (-1,s[-1]))
    loss = K.sum(keras.losses.categorical_crossentropy(y_true, y_pred))
    num = K.shape(y_true)[0]
    num=tf.cast(num,tf.float32)
    return K.mean(loss)/num
```

```
def metric1(y_true, y_pred):
    s = K.shape(y_pred)

# reshape such that w and h dim are multiplied together
    y_true_reshaped = K.reshape( y_true, (-1,s[-1]) )
    y_pred_reshaped = K.reshape( y_pred, (-1, s[-1]) )

# correctly classified
    clf_pred = K.argmax(y_pred_reshaped,axis = -1)
    y_true = K.argmax(y_true_reshaped,axis = -1)
    correct_pixels_per_class = K_cast( K_eggal(clf_pred_v_true) = dtype='float32')  #if_eggal
```

```
COTTECT_PITCETS_PET_CIASS - W.CASC! W.EAMAI(CIT_PIECE, N_CIRC!, ACTAC- ITOACSS ) #IT EAMAI
    return K.sum(correct pixels per class) / K.cast(K.prod(s[:-1]), dtype='float32') #accuracy
In [0]:
def metric2(y_true, y_pred):
    s = K.shape(y pred)
    # correctly classified
    clf pred = K.argmax(y pred, axis = -1)
    y true = K.argmax(y true,axis = -1)
    correct pixels per class = K.cast(K.all( K.equal(clf pred,y true),axis=-1), dtype='float32') #i
f equal
    return K.sum(correct_pixels_per_class) / K.cast(K.prod(s[0]), dtype='float32') #accuracy
In [0]:
model.compile(loss = custom loss,optimizer='adam',metrics=[metric1,metric2])
In [0]:
model.load weights("idfy1.h5")
Training
In [0]:
datagen = ImageDataGenerator(width shift range=0.14,
                                  height shift range=0.08,
                                  fill mode='constant',
                                  zoom range = 0.1,
                                  rotation_range = 10,
                                  #rescale =1./255
mcp_save = ModelCheckpoint('idfy1.h5', save_best_only=True, monitor='val_loss', mode='min', verbose=
def scheduler(epoch):
    if epoch <3 :</pre>
        return 0.001/5
    elif epoch < 10:</pre>
       return 0.001/10
    elif epoch < 15:</pre>
       return 0.00001
    elif epoch <30:</pre>
       return 0.00001/2
n = X_train.shape[0]
lr reduce = LearningRateScheduler(scheduler, verbose = 1)
history = model.fit_generator(datagen.flow(X_train, y_train,batch_size=64),
                          epochs = 30,
                          steps per epoch=n//64,
                          callbacks=[lr reduce,mcp save],
                          validation data=(X test, y test))
Epoch 1/30
Epoch 00001: LearningRateScheduler setting learning rate to 0.0002.
```

8/8 [=========] - 12s 2s/step - loss: 0.5540 - metric1: 0.9953 - metric2: 0.9
705 - val_loss: 1.0427 - val_metric1: 0.9232 - val_metric2: 0.7448

Epoch 00001: val_loss improved from inf to 1.04265, saving model to idfy1.h5

Epoch 2/30

Epoch 00002: LearningRateScheduler setting learning rate to 0.0002.

```
8/8 [============] - 13s 2s/step - loss: 0.5444 - metricl: 0.9980 - metric2: 0.9
844 - val loss: 1.0200 - val metric1: 0.9034 - val metric2: 0.6128
Epoch 00002: val loss improved from 1.04265 to 1.02004, saving model to idfy1.h5
Epoch 3/30
Epoch 00003: LearningRateScheduler setting learning rate to 0.0002.
648 - val loss: 1.0594 - val metric1: 0.9147 - val metric2: 0.7031
Epoch 00003: val loss did not improve from 1.02004
Epoch 4/30
Epoch 00004: LearningRateScheduler setting learning rate to 0.0001.
724 - val loss: 1.0557 - val metric1: 0.9089 - val metric2: 0.7083
Epoch 00004: val loss did not improve from 1.02004
Epoch 5/30
Epoch 00005: LearningRateScheduler setting learning rate to 0.0001.
8/8 [===========] - 11s 1s/step - loss: 0.5984 - metric1: 0.9919 - metric2: 0.9
388 - val loss: 1.0583 - val metric1: 0.8956 - val metric2: 0.5816
Epoch 00005: val loss did not improve from 1.02004
Epoch 6/30
Epoch 00006: LearningRateScheduler setting learning rate to 0.0001.
646 - val loss: 1.0796 - val metric1: 0.9021 - val metric2: 0.5868
Epoch 00006: val loss did not improve from 1.02004
Epoch 7/30
Epoch 00007: LearningRateScheduler setting learning rate to 0.0001.
824 - val loss: 0.9968 - val metric1: 0.9219 - val metric2: 0.7396
Epoch 00007: val loss improved from 1.02004 to 0.99679, saving model to idfy1.h5
Epoch 8/30
Epoch 00008: LearningRateScheduler setting learning rate to 0.0001.
8/8 [===========] - 12s 2s/step - loss: 0.5390 - metric1: 0.9990 - metric2: 0.9
922 - val loss: 1.0066 - val metric1: 0.9060 - val metric2: 0.6128
Epoch 00008: val loss did not improve from 0.99679
Epoch 9/30
Epoch 00009: LearningRateScheduler setting learning rate to 0.0001.
8/8 [===========] - 12s 1s/step - loss: 0.5338 - metric1: 0.9988 - metric2: 0.9
922 - val_loss: 1.0025 - val_metric1: 0.9093 - val_metric2: 0.6285
Epoch 00009: val loss did not improve from 0.99679
Epoch 10/30
Epoch 00010: LearningRateScheduler setting learning rate to 0.0001.
8/8 [===========] - 12s 1s/step - loss: 0.5339 - metric1: 0.9968 - metric2: 0.9
744 - val loss: 1.0216 - val metric1: 0.9212 - val metric2: 0.7344
Epoch 00010: val loss did not improve from 0.99679
Epoch 11/30
Epoch 00011: LearningRateScheduler setting learning rate to 1e-05.
8/8 [============ ] - 12s 1s/step - loss: 0.5480 - metric1: 0.9951 - metric2: 0.9
646 - val loss: 0.9872 - val metric1: 0.9212 - val metric2: 0.7240
Epoch 00011: val loss improved from 0.99679 to 0.98720, saving model to idfy1.h5
Epoch 12/30
Epoch 00012: LearningRateScheduler setting learning rate to 1e-05.
941 - val loss: 0.9816 - val metric1: 0.8974 - val metric2: 0.6285
Epoch 00012: val_loss improved from 0.98720 to 0.98162, saving model to idfy1.h5
Epoch 13/30
Epoch 00013: LearningRateScheduler setting learning rate to 1e-05.
```

```
8/8 [============= ] - 12s 1s/step - loss: 0.5395 - metric1: 0.9948 - metric2: 0.9
744 - val loss: 0.9794 - val metric1: 0.9106 - val metric2: 0.6285
Epoch 00013: val loss improved from 0.98162 to 0.97943, saving model to idfy1.h5
Epoch 14/30
Epoch 00014: LearningRateScheduler setting learning rate to 1e-05.
763 - val_loss: 0.9696 - val_metric1: 0.9271 - val_metric2: 0.7552
Epoch 00014: val loss improved from 0.97943 to 0.96959, saving model to idfy1.h5
Epoch 15/30
Epoch 00015: LearningRateScheduler setting learning rate to 1e-05.
941 - val loss: 0.9579 - val metric1: 0.9138 - val metric2: 0.6493
Epoch 00015: val loss improved from 0.96959 to 0.95785, saving model to idfy1.h5
Epoch 16/30
Epoch 00016: LearningRateScheduler setting learning rate to 5e-06.
941 - val_loss: 0.9514 - val_metric1: 0.9132 - val_metric2: 0.6441
Epoch 00016: val_loss improved from 0.95785 to 0.95136, saving model to idfy1.h5
Epoch 17/30
Epoch 00017: LearningRateScheduler setting learning rate to 5e-06.
783 - val loss: 0.9468 - val metric1: 0.9138 - val metric2: 0.6441
Epoch 00017: val loss improved from 0.95136 to 0.94681, saving model to idfy1.h5
Epoch 18/30
Epoch 00018: LearningRateScheduler setting learning rate to 5e-06.
744 - val loss: 0.9498 - val metric1: 0.9138 - val metric2: 0.6493
Epoch 00018: val_loss did not improve from 0.94681
Epoch 19/30
Epoch 00019: LearningRateScheduler setting learning rate to 5e-06.
980 - val loss: 0.9514 - val metric1: 0.9138 - val metric2: 0.6493
Epoch 00019: val loss did not improve from 0.94681
Epoch 20/30
Epoch 00020: LearningRateScheduler setting learning rate to 5e-06.
8/8 [===========] - 12s 1s/step - loss: 0.5057 - metric1: 0.9988 - metric2: 0.9
922 - val loss: 0.9515 - val metric1: 0.9132 - val metric2: 0.6493
Epoch 00020: val loss did not improve from 0.94681
Epoch 21/30
Epoch 00021: LearningRateScheduler setting learning rate to 5e-06.
941 - val_loss: 0.9485 - val_metric1: 0.9264 - val_metric2: 0.7604
Epoch 00021: val loss did not improve from 0.94681
Epoch 22/30
Epoch 00022: LearningRateScheduler setting learning rate to 5e-06.
961 - val loss: 0.9463 - val metric1: 0.9271 - val metric2: 0.7604
Epoch 00022: val loss improved from 0.94681 to 0.94630, saving model to idfy1.h5
Epoch 23/30
Epoch 00023: LearningRateScheduler setting learning rate to 5e-06.
922 - val_loss: 0.9433 - val_metric1: 0.9264 - val_metric2: 0.7552
Epoch 00023: val loss improved from 0.94630 to 0.94326, saving model to idfy1.h5
Epoch 24/30
Epoch 00024: LearningRateScheduler setting learning rate to 5e-06.
```

```
8/8 [===========] - 13s 2s/step - loss: 0.4999 - metric1: 0.9988 - metric2: 0.9
902 - val loss: 0.9423 - val_metric1: 0.9132 - val_metric2: 0.6441
Epoch 00024: val loss improved from 0.94326 to 0.94227, saving model to idfy1.h5
Epoch 25/30
Epoch 00025: LearningRateScheduler setting learning rate to 5e-06.
961 - val_loss: 0.9408 - val_metric1: 0.9132 - val_metric2: 0.6441
Epoch 00025: val loss improved from 0.94227 to 0.94077, saving model to idfy1.h5
Epoch 26/30
Epoch 00026: LearningRateScheduler setting learning rate to 5e-06.
8/8 [========= ] - 12s 1s/step - loss: 0.5247 - metric1: 0.9983 - metric2: 0.9
883 - val loss: 0.9411 - val metric1: 0.9138 - val metric2: 0.6441
Epoch 00026: val loss did not improve from 0.94077
Epoch 27/30
Epoch 00027: LearningRateScheduler setting learning rate to 5e-06.
8/8 [===========] - 12s 1s/step - loss: 0.5006 - metric1: 0.9993 - metric2: 0.9
941 - val_loss: 0.9423 - val_metric1: 0.9284 - val_metric2: 0.7552
Epoch 00027: val loss did not improve from 0.94077
Epoch 28/30
Epoch 00028: LearningRateScheduler setting learning rate to 5e-06.
961 - val loss: 0.9452 - val metric1: 0.9284 - val metric2: 0.7552
Epoch 00028: val loss did not improve from 0.94077
Epoch 29/30
Epoch 00029: LearningRateScheduler setting learning rate to 5e-06.
763 - val_loss: 0.9452 - val_metric1: 0.9290 - val_metric2: 0.7604
Epoch 00029: val_loss did not improve from 0.94077
Epoch 30/30
Epoch 00030: LearningRateScheduler setting learning rate to 5e-06.
766 - val loss: 0.9445 - val metric1: 0.9284 - val metric2: 0.7604
Epoch 00030: val loss did not improve from 0.94077
```

Utility funtion for plotting

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
In [0]:
```

```
nb_epoch=30
```

```
In [49]:
```

```
%matplotlib inline
```

```
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch'); ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test accuracy: 0.9041666984558105

Test Accuracy: 90%

Let's test on test data if it works good !!

```
In [39]:
```

```
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import random
n = X_test.shape[0]
count = 0
pred = model.predict(X_test)
pred = np.argmax(pred,axis=-1)
true = np.argmax(y_test,axis = -1)
for i in range(n):
    if np.all(pred[i,:]==true[i]):
        count +=1
print(count/n)
```

0.6335877862595419

It gave me accurate results upto 63.4%. Thats great for this small dataset

In [40]:

```
k = random.randint(0,100)
imgplot = plt.imshow(X_test[k,:,:].reshape(64,256))
arr = []
for i in range(8):
    arr.append(dic[pred[k,i]])
print(arr)
```

```
['5', 'B', '2', '9', '6', '7', '0', ' ']
```



In [42]:

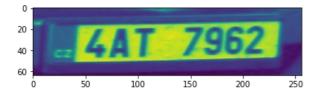
```
k = random.randint(0,100)
imgplot = plt.imshow(X_test[k,:,:].reshape(64,256))
arr = []
for i in range(8):
    arr.append(dic[pred[k,i]])
print(arr)
['7'. 'B'. '1'. '4'. '1'. '5'. '6'. ' ']
```

```
0
20
40
60
50
100
150
200
250
```

In [41]:

```
k = random.randint(0,100)
imgplot = plt.imshow(X_test[k,:,:,:].reshape(64,256))
arr = []
for i in range(8):
    arr.append(dic[pred[k,i]])
print(arr)
```

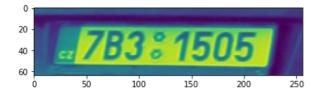
```
['4', 'A', 'X', '7', '9', '6', '2', ' ']
```



In [44]:

```
k = random.randint(0,100)
imgplot = plt.imshow(X_test[k,:,:].reshape(64,256))
arr = []
for i in range(8):
    arr.append(dic[pred[k,i]])
print(arr)
```

```
['7', 'B', '3', '1', '5', '0', '5', ' ']
```



In [51]:

```
k = random.randint(0,100)
imgplot = plt.imshow(X_test[k,:,:].reshape(64,256))
arr = []
for i in range(8):
    arr.append(dic[pred[k,i]])
print(arr)
```

```
['9', 'B', '4', '6', '9', '8', '8', ' ']
```



In [53]:

```
imgplot = plt.imshow(X_test[k,:,:].reshape(64,256))
arr = []
for i in range(8):
    arr.append(dic[pred[k,i]])
print(arr)
```

```
['7', 'B', '0', '8', '1', '9', '9', ' ']
```



Conclusion:

- Test Accuracy 90%
- Even this first cut solution gave tremendous results.
- The dataset was very small so, it had very less training data .
- Need to modify CNN architecture and need more data for better accuracy.
- Modifications : YOLO V3 for text detection would have given better results as per different research papers.